

Languages as Hyperplanes

Grammatical Inference with String Kernels

Alex Clark Christophe Costa Florencio Chris Watkins

Department of Computer Science
Royal Holloway, University of London

European Conference on Machine Learning, 2006

Acknowledgement of Support

We would like to acknowledge support from the Pascal Network
of Excellence,

in the form of a 'pump-priming' grant 2005-2006 for
Grammatical Inference with String Kernels

Acknowledgement of Support

We would like to acknowledge support from the Pascal Network
of Excellence,
in the form of a ‘pump-priming’ grant 2005-2006 for
Grammatical Inference with String Kernels

Original Motivation

Language Learning from Positive Examples

Learning a language from positive examples only

- ▶ Children learn complex grammars
 - ▶ without being given negative examples
 - ▶ without being corrected
- ▶ Classic problem in computational linguistics
- ▶ Useful formulation in practice because informative negative examples difficult to obtain/generate

Learnability criterion:

- ▶ Learner is presented with a sequence of sentences from a language
- ▶ After some (preferably small) number of sentences, learner acquires exact description of language

Formalising Language Learning

- ▶ Finite alphabet Σ
- ▶ Set of finite sequences Σ^*
- ▶ A *language* L is a set of finite sequences, $L \subseteq \Sigma^*$
- ▶ Problem: learn to recognise sequences in L from presentation of positive examples only.

After finite number of examples, learner should acquire **exact** description of language.

PAC-learning type criteria also possible...

Formalising Language Learning

- ▶ Finite alphabet Σ
- ▶ Set of finite sequences Σ^*
- ▶ A *language* L is a set of finite sequences, $L \subseteq \Sigma^*$
- ▶ Problem: learn to recognise sequences in L from presentation of positive examples only.

After finite number of examples, learner should acquire **exact** description of language.

PAC-learning type criteria also possible...

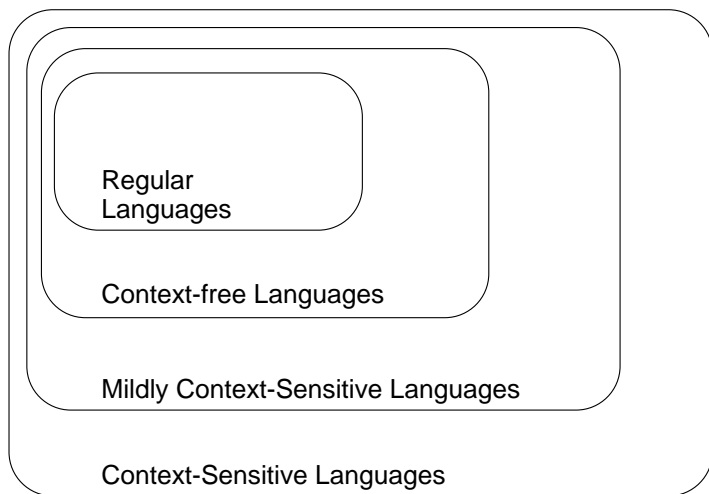
Formalising Language Learning

- ▶ Finite alphabet Σ
- ▶ Set of finite sequences Σ^*
- ▶ A *language* L is a set of finite sequences, $L \subseteq \Sigma^*$
- ▶ Problem: learn to recognise sequences in L from presentation of positive examples only.

After finite number of examples, learner should acquire **exact** description of language.

PAC-learning type criteria also possible...

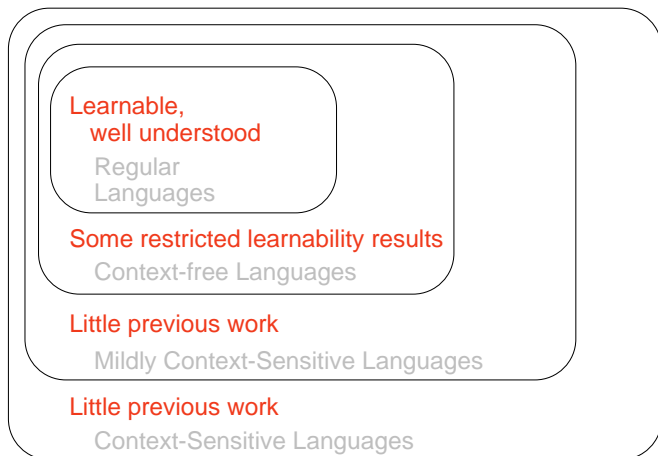
Representing Languages: the Chomsky Hierarchy



Grammars of increasing intractability

Natural language believed to be mostly context-free, but some mildly context-sensitive phenomena.

Grammar Induction in Computational Linguistics



But Chomsky hierarchy does not match intuitive notions of complexity...

Some intuitively simple languages are MCS or CS: how to learn them?

Planar Languages: a New Approach

Map strings to points in a Euclidean feature space $\phi : \Sigma^* \mapsto H$

To define a language, specify a region U in feature space

Language is all strings s such that $\phi(s) \in U$

Language is *pre-image* of U in Σ^*

So, two questions:

What feature spaces?

What type of region?

Planar Languages: a New Approach

Map strings to points in a Euclidean feature space $\phi : \Sigma^* \mapsto H$

To define a language, specify a region U in feature space

Language is all strings s such that $\phi(s) \in U$

Language is *pre-image* of U in Σ^*

So, two questions:

What feature spaces?

What type of region?

Planar Languages: a New Approach

Map strings to points in a Euclidean feature space $\phi : \Sigma^* \mapsto H$

To define a language, specify a region U in feature space

Language is all strings s such that $\phi(s) \in U$

Language is *pre-image* of U in Σ^*

So, two questions:

What feature spaces?

What type of region?

What Type of Region?

Hmm....could use half-spaces? Conventional... Spheres?
Manifolds? Tricky... All of these are possible...But

Hyperplanes seem natural
Finite-dimensional hyperplanes are particularly nice...let's see
why

What Type of Region?

Hmm....could use half-spaces? Conventional... Spheres?
Manifolds? Tricky... All of these are possible...But

Hyperplanes seem natural
Finite-dimensional hyperplanes are particularly nice...let's see
why

What Type of Region?

Hmm....could use half-spaces? Conventional... Spheres?

Manifolds? Tricky... All of these are possible...But

Hyperplanes seem natural

Finite-dimensional hyperplanes are particularly nice...let's see why

What Type of Region?

Hmm....could use half-spaces? Conventional... Spheres?
Manifolds? Tricky... All of these are possible...But

Hyperplanes seem natural

Finite-dimensional hyperplanes are particularly nice...let's see
why

What Type of Region?

Hmm....could use half-spaces? Conventional... Spheres?
Manifolds? Tricky... All of these are possible...But

Hyperplanes seem natural

Finite-dimensional hyperplanes are particularly nice...let's see
why

What Type of Region?

Hmm....could use half-spaces? Conventional... Spheres?
Manifolds? Tricky... All of these are possible...But

Hyperplanes seem natural

Finite-dimensional hyperplanes are particularly nice...let's see
why

What Type of Region?

Hmm....could use half-spaces? Conventional... Spheres?
Manifolds? Tricky... All of these are possible...But

Hyperplanes seem natural

Finite-dimensional hyperplanes are particularly nice...let's see
why

What Type of Region?

Hmm....could use half-spaces? Conventional... Spheres?
Manifolds? Tricky... All of these are possible...But

Hyperplanes seem natural

Finite-dimensional hyperplanes are particularly nice...let's see why

What Type of Region?

Hmm....could use half-spaces? Conventional... Spheres?
Manifolds? Tricky... All of these are possible...But

Hyperplanes seem natural
Finite-dimensional hyperplanes are particularly nice...let's see
why

Planar Languages

Definition: A language L is *phi-planar* if

$$L = \{w : \phi(w) \in U \subseteq H\}$$

where U is a r -dimensional hyperplane (not necessarily containing the origin)

Useful definition because:

- ▶ hyperplane defined by *linear equality constraints* in feature space, often easy to interpret
- ▶ hyperplane learned exactly when $r + 1$ linearly independent strings observed
- ▶ very simple learning algorithm

Planar Languages

Definition: A language L is *phi-planar* if

$$L = \{w : \phi(w) \in U \subseteq H\}$$

where U is a r -dimensional hyperplane (not necessarily containing the origin)

Useful definition because:

- ▶ hyperplane defined by linear equality constraints in feature space, often easy to interpret
- ▶ hyperplane learned *exactly* when $r + 1$ linearly independent strings observed
- ▶ very simple learning algorithm

Planar Languages

Definition: A language L is *phi-planar* if

$$L = \{w : \phi(w) \in U \subseteq H\}$$

where U is a r -dimensional hyperplane (not necessarily containing the origin)

Useful definition because:

- ▶ hyperplane defined by linear equality constraints in feature space, often easy to interpret
- ▶ hyperplane learned exactly when $r + 1$ linearly independent strings observed
- ▶ very simple learning algorithm

Planar Languages

Definition: A language L is *phi-planar* if

$$L = \{w : \phi(w) \in U \subseteq H\}$$

where U is a r -dimensional hyperplane (not necessarily containing the origin)

Useful definition because:

- ▶ hyperplane defined by linear equality constraints in feature space, often easy to interpret
- ▶ hyperplane learned exactly when $r + 1$ linearly independent strings observed
- ▶ very simple learning algorithm

Example: Strings with equal numbers of a's and b's


Let $\Sigma = \{a, b\}$

Consider $L = \{s \in \Sigma^* : |s|_a = |s|_b\}$
where $|s|_a$ is the number of a 's in s

i.e. L consists of strings with equal numbers of a and b

Consider feature space H with two dimensions: $|s|_a$ and $|s|_b$

Example: Strings with equal numbers of a's and b's



Hyperplane is a line

Language acquired once 2 different points on line found

Learning Hyperplanes: KPCA of centred Gram Matrix

Algorithm (for kernel-defined feature space mapping):

- ▶ Form Gram matrix for positive data
- ▶ Centre Gram matrix (record the displacement necessary)
- ▶ Compute principal components of Gram matrix. Keep significant components.
- ▶ Hyperplane now defined!

Convergence

Planar Target Language:

- ▶ Convergence is rapid and exact
- ▶ Key factor is rank r of hyperplane
- ▶ At least r independent examples required to learn language
- ▶ Incomplete learning produces *false negatives* in test set

Non-Planar Target Language:

- ▶ Convergence to lowest-dimensional hyperplane containing language which may be all of Σ^*
- ▶ Learning non-planar language produces *false positives* in test set

Convergence

Planar Target Language:

- ▶ Convergence is rapid and exact
- ▶ Key factor is rank r of hyperplane
- ▶ At least r independent examples required to learn language
- ▶ Incomplete learning produces *false negatives* in test set

Non-Planar Target Language:

- ▶ Convergence to lowest-dimensional hyperplane containing language which may be all of Σ^*
- ▶ Learning non-planar language produces *false positives* in test set

Convergence

Planar Target Language:

- ▶ Convergence is rapid and exact
- ▶ Key factor is rank r of hyperplane
- ▶ At least r independent examples required to learn language
- ▶ Incomplete learning produces *false negatives* in test set

Non-Planar Target Language:

- ▶ Convergence to lowest-dimensional hyperplane containing language which may be all of Σ^*
- ▶ Learning non-planar language produces *false positives* in test set

Convergence

Planar Target Language:

- ▶ Convergence is rapid and exact
- ▶ Key factor is rank r of hyperplane
- ▶ At least r independent examples required to learn language
- ▶ Incomplete learning produces *false negatives* in test set

Non-Planar Target Language:

- ▶ Convergence to lowest-dimensional hyperplane containing language which may be all of Σ^*
- ▶ Learning non-planar language produces *false positives* in test set

Convergence

Planar Target Language:

- ▶ Convergence is rapid and exact
- ▶ Key factor is rank r of hyperplane
- ▶ At least r independent examples required to learn language
- ▶ Incomplete learning produces *false negatives* in test set

Non-Planar Target Language:

- ▶ Convergence to lowest-dimensional hyperplane containing language which may be all of Σ^*
- ▶ Learning non-planar language produces *false positives* in test set

Convergence

Planar Target Language:

- ▶ Convergence is rapid and exact
- ▶ Key factor is rank r of hyperplane
- ▶ At least r independent examples required to learn language
- ▶ Incomplete learning produces *false negatives* in test set

Non-Planar Target Language:

- ▶ Convergence to lowest-dimensional hyperplane containing language which may be all of Σ^*
- ▶ Learning non-planar language produces *false positives* in test set

Convergence

Planar Target Language:

- ▶ Convergence is rapid and exact
- ▶ Key factor is rank r of hyperplane
- ▶ At least r independent examples required to learn language
- ▶ Incomplete learning produces *false negatives* in test set

Non-Planar Target Language:

- ▶ Convergence to lowest-dimensional hyperplane containing language which may be all of Σ^*
- ▶ Learning non-planar language produces *false positives* in test set

Feature Space: Subsequence Counts

We consider two feature spaces (can be induced by string-kernels):

- ▶ Sub-sequence counts
- ▶ Gap-weighted sub-sequence counts

Example:

Let $s = ababab$

Count of ab in s is $|s|_{ab} = 6$

Gap-weighted count of ab in s with weighting factor $\frac{1}{2}$ is

$$3 + 2 \times \frac{1}{4} + 1 \times \frac{1}{64}$$

Feature Space: Subsequence Counts

We consider two feature spaces (can be induced by string-kernels):

- ▶ Sub-sequence counts
- ▶ Gap-weighted sub-sequence counts

Example:

Let $s = ababab$

Count of ab in s is $|s|_{ab} = 6$

Gap-weighted count of ab in s with weighting factor $\frac{1}{2}$ is

$$3 + 2 \times \frac{1}{4} + 1 \times \frac{1}{64}$$

Feature Space: Subsequence Counts

We consider two feature spaces (can be induced by string-kernels):

- ▶ Sub-sequence counts
- ▶ Gap-weighted sub-sequence counts

Example:

Let $s = ababab$

Count of ab in s is $|s|_{ab} = 6$

Gap-weighted count of ab in s with weighting factor $\frac{1}{2}$ is

$$3 + 2 \times \frac{1}{4} + 1 \times \frac{1}{64}$$

Feature Space: Subsequence Counts

We consider two feature spaces (can be induced by string-kernels):

- ▶ Sub-sequence counts
- ▶ Gap-weighted sub-sequence counts

Example:

Let $s = ababab$

Count of ab in s is $|s|_{ab} = 6$

Gap-weighted count of ab in s with weighting factor $\frac{1}{2}$ is

$$3 + 2 \times \frac{1}{4} + 1 \times \frac{1}{64}$$

Experiments: Setup

- ▶ A range of languages tried, mostly already defined in computational linguistics literature
- ▶ 500 positive examples given; tested on 500 positive and negative examples
- ▶ specially informative negative test examples generated where necessary

We use two feature spaces:

- ▶ Counts of sub-sequences of lengths **one and two**
- ▶ Gap-weighted counts of sub-sequences of lengths **one and two**

Experiments: Setup

- ▶ A range of languages tried, mostly already defined in computational linguistics literature
- ▶ 500 positive examples given; tested on 500 positive and negative examples
- ▶ specially informative negative test examples generated where necessary

We use two feature spaces:

- ▶ Counts of sub-sequences of lengths **one and two**
- ▶ Gap-weighted counts of sub-sequences of lengths **one and two**

Experiments: Even

Even
(Regular)

Even number of symbols
Alphabet $\{a, b, c\}$

abcb, ba,
babacc,
aaaa

Bracket
(CF)

Balanced brackets
Alphabet $\{(,)\}$

(), $()()$,
 $((()))$

	PCFG		HMM		SUBS			GPWT		
	FP	FN	FP	FN	FP	FN	R	FP	FN	R
Even	0	0	0	0	100	0	12	100	0	12
Bracket	0	0	3.4	1.3	10.8	0	3	10.8	0	5

Experiments: Even

Even
(Regular) Even number of symbols
Alphabet $\{a, b, c\}$ abcb, ba,
babacc,
aaaa

Bracket
(CF) Balanced brackets
Alphabet $\{(,)\}$ (), ()(),
(()(()))

	PCFG		HMM		SUBS		GPWT			
	FP	FN	FP	FN	FP	FN	R	FP	FN	R
Even	0	0	0	0	100	0	12	100	0	12
Bracket	0	0	3.4	1.3	10.8	0	3	10.8	0	5

A Regular Language with Long-Range Dependency

GWLang
(Regular)

Strings of form $uavbw$,
with
 $u, v, w \in \{c, d, e, f, g, h\}^*$,
with $|v| = 3$

acccb,
feaghfbec,
gahefbdg

	PCFG		HMM		SUBS			GPWT		
	FP	FN	FP	FN	FP	FN	R	FP	FN	R
GWLang	-	-	25.4	5	0	0	53	0	1	64

Planar Languages not Learned by HMMs or PCFGs

A German Verb Construction (Mildly Context Sensitive)

Strings in two halves: first half from $\{a, b, c, d\}$, second from $\{e, f, g, h\}$

For each a in first half, there must be a e in the second half, etc, but in any order

Example strings: *abfe*, *abcdegfh*, *aabefe*

Planar Languages not Learned by HMMs or PCFGs

Strings of the form:

n characters from $\{a, b\}$,

followed by n characters from $\{c, d\}$,

followed by n characters from $\{e, f\}$

etc

babcccfef, adf, aaaaccccefef

Depending on the number of sections these languages are context-free, mildly context-sensitive, or context-sensitive. Easily defined by linear equality constraints.

Languages with Long-Range Sequence Dependencies

For example:

Palindromes: *abcabbaabbacba*

Single Repeats: *abcaabca*, *abbababbab*

Some of these languages context-free; others mildly context-sensitive

These are **approximately planar** in the sense that, for a given length of subsequences in the feature space, only sequences up to a certain length are fully specified by the feature vector.

e.g. *abba* and *baab* have the same feature vector for subsequences of length up to 2

For subsequences of length 3, shortest pair of indistinguishable sequences is of length 7.

Languages with Long-Range Sequence Dependencies

For example:

Palindromes: *abcabbaabbacba*

Single Repeats: *abcaabca*, *abbababbab*

Some of these languages context-free; others mildly context-sensitive

These are **approximately planar** in the sense that, for a given length of subsequences in the feature space, only sequences up to a certain length are fully specified by the feature vector.

e.g. *abba* and *baab* have the same feature vector for subsequences of length up to 2

For subsequences of length 3, shortest pair of indistinguishable sequences is of length 7.

Languages with Long-Range Sequence Dependencies

For example:

Palindromes: *abcabbaabbacba*

Single Repeats: *abcaabca*, *abbababbab*

Some of these languages context-free; others mildly context-sensitive

These are **approximately planar** in the sense that, for a given length of subsequences in the feature space, only sequences up to a certain length are fully specified by the feature vector.

e.g. *abba* and *baab* have the same feature vector for subsequences of length up to 2

For subsequences of length 3, shortest pair of indistinguishable sequences is of length 7.

Languages with Long-Range Sequence Dependencies

For example:

Palindromes: *abcabbaabbacba*

Single Repeats: *abcaabca*, *abbababbab*

Some of these languages context-free; others mildly context-sensitive

These are **approximately planar** in the sense that, for a given length of subsequences in the feature space, only sequences up to a certain length are fully specified by the feature vector.

e.g. *abba* and *baab* have the same feature vector for subsequences of length up to 2

For subsequences of length 3, shortest pair of indistinguishable sequences is of length 7.

Languages with Long-Range Sequence Dependencies

For example:

Palindromes: *abcabbaabbacba*

Single Repeats: *abcaabca*, *abbababbab*

Some of these languages context-free; others mildly context-sensitive

These are **approximately planar** in the sense that, for a given length of subsequences in the feature space, only sequences up to a certain length are fully specified by the feature vector.
e.g. *abba* and *baab* have the same feature vector for subsequences of length up to 2

For subsequences of length 3, shortest pair of indistinguishable sequences is of length 7.

Languages with Long-Range Sequence Dependencies

For example:

Palindromes: *abcabbaabbacba*

Single Repeats: *abcaabca*, *abbababbab*

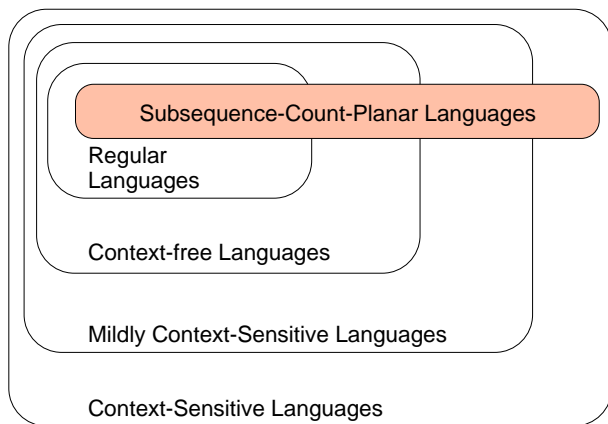
Some of these languages context-free; others mildly context-sensitive

These are **approximately planar** in the sense that, for a given length of subsequences in the feature space, only sequences up to a certain length are fully specified by the feature vector.

e.g. *abba* and *baab* have the same feature vector for subsequences of length up to 2

For subsequences of length 3, shortest pair of indistinguishable sequences is of length 7.

Subsequence Planar Languages Cross-Cut Chomsky Hierarchy



Summary

- ▶ Planar languages: a new(?) approach to language induction from positive examples
- ▶ Natural to represent some languages in terms of linear constraints on sub-sequences
- ▶ Certain types of long-range dependency of intervals, counts, and sequences can be
 - ▶ neatly represented
 - ▶ reliably learned
- ▶ For simple subsequence kernels, planar languages cross-cut Chomsky hierarchy
- ▶ Our example languages are intuitively simple
Intuitively, earning them should be simple
Our approach makes it simple
- ▶ Further directions
 - ▶ Characterise set of planar languages with respect to various kernels
 - ▶ String processing in features space...

Summary

- ▶ Planar languages: a new(?) approach to language induction from positive examples
- ▶ Natural to represent some languages in terms of linear constraints on sub-sequences
- ▶ Certain types of long-range dependency of intervals, counts, and sequences can be
 - ▶ neatly represented
 - ▶ reliably learned
- ▶ For simple subsequence kernels, planar languages cross-cut Chomsky hierarchy
- ▶ Our example languages are intuitively simple
Intuitively, earning them should be simple
Our approach makes it simple
- ▶ Further directions
 - ▶ Characterise set of planar languages with respect to various kernels
 - ▶ String processing in features space...

Summary

- ▶ Planar languages: a new(?) approach to language induction from positive examples
- ▶ Natural to represent some languages in terms of linear constraints on sub-sequences
- ▶ Certain types of long-range dependency of intervals, counts, and sequences can be
 - ▶ neatly represented
 - ▶ reliably learned
- ▶ For simple subsequence kernels, planar languages cross-cut Chomsky hierarchy
- ▶ Our example languages are intuitively simple
Intuitively, earning them should be simple
Our approach makes it simple
- ▶ Further directions
 - ▶ Characterise set of planar languages with respect to various kernels
 - ▶ String processing in features space...

Summary

- ▶ Planar languages: a new(?) approach to language induction from positive examples
- ▶ Natural to represent some languages in terms of linear constraints on sub-sequences
- ▶ Certain types of long-range dependency of intervals, counts, and sequences can be
 - ▶ neatly represented
 - ▶ reliably learned
- ▶ For simple subsequence kernels, planar languages cross-cut Chomsky hierarchy
- ▶ Our example languages are intuitively simple
Intuitively, earning them should be simple
Our approach makes it simple
- ▶ Further directions
 - ▶ Characterise set of planar languages with respect to various kernels
 - ▶ String processing in features space...

Summary

- ▶ Planar languages: a new(?) approach to language induction from positive examples
- ▶ Natural to represent some languages in terms of linear constraints on sub-sequences
- ▶ Certain types of long-range dependency of intervals, counts, and sequences can be
 - ▶ neatly represented
 - ▶ reliably learned
- ▶ For simple subsequence kernels, planar languages cross-cut Chomsky hierarchy
- ▶ Our example languages are intuitively simple
Intuitively, learning them should be simple
Our approach makes it simple
- ▶ Further directions
 - ▶ Characterise set of planar languages with respect to various kernels
 - ▶ String processing in features space...

Summary

- ▶ Planar languages: a new(?) approach to language induction from positive examples
- ▶ Natural to represent some languages in terms of linear constraints on sub-sequences
- ▶ Certain types of long-range dependency of intervals, counts, and sequences can be
 - ▶ neatly represented
 - ▶ reliably learned
- ▶ For simple subsequence kernels, planar languages cross-cut Chomsky hierarchy
- ▶ Our example languages are intuitively simple
Intuitively, earning them should be simple
Our approach makes it simple
- ▶ Further directions
 - ▶ Characterise set of planar languages with respect to various kernels
 - ▶ String processing in features space...

Summary

- ▶ Planar languages: a new(?) approach to language induction from positive examples
- ▶ Natural to represent some languages in terms of linear constraints on sub-sequences
- ▶ Certain types of long-range dependency of intervals, counts, and sequences can be
 - ▶ neatly represented
 - ▶ reliably learned
- ▶ For simple subsequence kernels, planar languages cross-cut Chomsky hierarchy
- ▶ Our example languages are intuitively simple
Intuitively, learning them should be simple
Our approach makes it simple
- ▶ Further directions
 - ▶ Characterise set of planar languages with respect to various kernels
 - ▶ String processing in features space...

Summary

- ▶ Planar languages: a new(?) approach to language induction from positive examples
- ▶ Natural to represent some languages in terms of linear constraints on sub-sequences
- ▶ Certain types of long-range dependency of intervals, counts, and sequences can be
 - ▶ neatly represented
 - ▶ reliably learned
- ▶ For simple subsequence kernels, planar languages cross-cut Chomsky hierarchy
- ▶ Our example languages are intuitively simple
Intuitively, learning them should be simple
Our approach makes it simple
- ▶ Further directions
 - ▶ Characterise set of planar languages with respect to various kernels
 - ▶ String processing in features space...